

Optimization and Computational Linear Algebra for Data Science

Final review problems

For review exercises on linear algebra, look at last year's final review exercises (available the course's website).

Problem 0.1. Let $A \in \mathbb{R}^{n \times m}$. Let $\sigma_1(A)$ be the largest singular value of A . Show that

$$\sigma_1(A) = \max_{\|x\|=1} \|Ax\|.$$

Problem 0.2. Let $A \in \mathbb{R}^{n \times m}$. Show that $A^T A$ and AA^T have the same non-zero eigenvalues.

Problem 0.3 (True or false?). For each of the following statement, say if they are true or false and justify your answer.

- For all $A \in \mathbb{R}^{n \times n}$, if λ is an eigenvalue of A then λ^2 is an eigenvalue of A^2 .
- For all $A \in \mathbb{R}^{n \times n}$, if σ is a singular value of A then σ^2 is a singular value of A^2 .
- For all symmetric matrix $A \in \mathbb{R}^{n \times n}$ the eigenvalues of A are singular values of A .

Problem 0.4. Let $A \in \mathbb{R}^{n \times m}$. Show that for all $u \in \text{Im}(A)$ and for all $v \in \text{Ker}(A^T)$ we have

$$\langle u, v \rangle = 0.$$

Problem 0.5. Let $A \in \mathbb{R}^{n \times m}$. Show that if A has linearly independent columns, then $A^\dagger = (A^T A)^{-1} A^T$.

Problem 0.6. Which of the following functions $f : \mathbb{R}^n \rightarrow \mathbb{R}$ are convex? Justify your answer

- $f(x) = \|x\|^2$.
- $f(x) = Ax$, for some $A \in \mathbb{R}^{n \times n}$.
- $f(x) = \sum_{i=1}^n x_i^3$.

Problem 0.7. Which of the following subset S of \mathbb{R}^n are convex? Justify your answer

- $S = \{x \in \mathbb{R}^n \mid \|x\|_\infty \leq 1\}$.
- $S = \{x \in \mathbb{R}^n \mid \|x\|_1 \geq 1\}$.
- $S = \{x \in \mathbb{R}^n \mid \|Ax\| < 1\}$, for some $A \in \mathbb{R}^{n \times n}$.

Problem 0.8. Show that we are performing PCA on n data points $a_1, \dots, a_n \in \mathbb{R}^d$ and keep only the first $k < d$ principal components of each point. We store the dimensionally reduced dataset in a $n \times k$ matrix B , where $B_{i,j}$ is the j^{th} principal component of the point a_i . Show that the columns of B are orthogonal.

Problem 0.9 (True or false?). For each of the following statement, say if they are true or false and justify your answer.

- If a continuous function $f : \mathbb{R} \rightarrow \mathbb{R}$ has a unique minimizer then f is convex.
- If a continuous function $f : \mathbb{R} \rightarrow \mathbb{R}$ is such that there exists x_0 such that f is decreasing on $(-\infty, x_0]$ and increasing on $[x_0, +\infty)$ then f is convex.
- A twice differentiable function $f : \mathbb{R} \rightarrow \mathbb{R}$ whose derivative f' is non-decreasing is convex.

Problem 0.10. Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a convex, differentiable function. Assume that there exist $x, y \in \mathbb{R}^n$ such that $\nabla f(x) = \nabla f(y) = 0$. Show that $\nabla f(\frac{1}{2}(x + y)) = 0$.

Problem 0.11. Assume that we are doing linear regression with the least-squares cost

$$f(x) = \|Ax - y\|^2$$

where $A \in \mathbb{R}^{n \times d}$ and $y \in \mathbb{R}^n$. Should you normalize the dataset A (that is, should we divide each column of A by its norm) to get better results (smaller training error or smaller test error on new data points)?

Suppose that we now want to use the lasso and minimize

$$f(x) = \|Ax - y\|^2 + \lambda \|x\|_1$$

for some $\lambda > 0$. Is there any reason why you might want to normalize the dataset in that case?

Problem 0.12. Compute the critical points of the following function and say if they are local minimizers, local maximizers or saddle points.

$$f(x, y, z) = x^2 + y^2 - z^2 \quad \text{and} \quad g(x, y) = 3x^2 + y^2 - 6x - 4y - 10.$$

Problem 0.13. Solve the following constrained minimization problem (find all the solutions to these problems).

1. Minimize $x + y + z$ subject to $e^{-x} + e^{-y} + e^{-z} = 1$.
2. Minimize $x^2 + y^2 + z^2$ subject to $xyz = 1$.

Problem 0.14. Assume that we are doing standard gradient descent to minimize the least-square cost

$$f(x) = \|Ax - y\|^2.$$

Assume that the columns of A are linearly dependent, meaning that $\text{Ker}(A) \neq \{0\}$. At which speed should gradient descent converge to the minimum? If now $\text{Ker}(A) = \{0\}$, at which speed should gradient descent converge? By speed, we only ask about the dependence in t , the number of iterations, of the gap $f(x_t) - \min f$, where x_t is the position of gradient descent after t iterations.

Problem 0.15. Let $A \in \mathbb{R}^{n \times d}$. Assume that the columns of A are linearly independent. How many steps of Newton's method do you need to minimize

$$\|Ax - y\|^2 ?$$

($y \in \mathbb{R}^n$ is a fixed vector). Justify your answer.

Problem 0.16. When running stochastic gradient descent, what are upsides and downsides of having a rapidly decaying learning rate?

