

Optimization and Computational Linear Algebra for Data Science

Lecture 7: Singular value decomposition

Léo MIOLANE · leo.miolane@gmail.com

October 25, 2019

Warning: *This material is not meant to be lecture notes. It only gathers the main concepts and results from the lecture, without any additional explanation, motivation, examples, figures...*

1 The Spectral Theorem

The main result of this section is the following “Spectral Theorem” which tells us that a symmetric matrix is diagonalizable in an orthonormal basis.

Theorem 1.1 (*Spectral Theorem*)

Let $A \in \mathbb{R}^{n \times n}$ be a **symmetric** matrix. Then there is a orthonormal basis of \mathbb{R}^n composed of eigenvectors of A .

Given an $n \times n$ symmetric matrix A , Theorem 1.1 tells us that one can find an orthonormal basis (v_1, \dots, v_n) of \mathbb{R}^n and scalars $\lambda_1, \dots, \lambda_n \in \mathbb{R}$ such that for all $i \in \{1, \dots, n\}$,

$$Av_i = \lambda_i v_i.$$

Let P be the $n \times n$ matrix whose columns are v_1, \dots, v_n . Since (v_1, \dots, v_n) is an orthonormal basis, we get that P is an orthogonal matrix. Let $D = \text{Diag}(\lambda_1, \dots, \lambda_n)$ and compute

$$AP = A \begin{pmatrix} | & | & \cdots & | \\ v_1 & v_2 & & v_n \\ | & | & & | \end{pmatrix} = \begin{pmatrix} | & | & \cdots & | \\ Av_1 & Av_2 & & Av_n \\ | & | & & | \end{pmatrix} = \begin{pmatrix} | & | & \cdots & | \\ \lambda_1 v_1 & \lambda_2 v_2 & & \lambda_n v_n \\ | & | & & | \end{pmatrix} = PD.$$

By multiplying by P^\top on both sides, we get $APP^\top = PDP^\top$. Recall now that P is orthogonal, therefore $PP^\top = \text{Id}_n$. We conclude that $A = PDP^\top$.

Theorem 1.2 (*Spectral Theorem, matrix formulation*)

Let $A \in \mathbb{R}^{n \times n}$ be a symmetric matrix. Then there exists an orthogonal matrix P and a diagonal matrix D of sizes $n \times n$, such that

$$A = PDP^\top.$$

Proposition 1.1

Let A be a $n \times n$ symmetric matrix and let $\lambda_1 \geq \dots \geq \lambda_n$ be its n eigenvalues and v_1, \dots, v_n be the associated orthonormal family of eigenvectors. Then

$$v_1 = \arg \max_{\|v\|=1} v^\top Av, \quad \text{and for } k = 2, \dots, n, \quad v_k = \arg \max_{\|v\|=1, v \perp v_1, \dots, v_{k-1}} v^\top Av.$$

Remark 1.1. Applying the proposition above to the matrix $-A$ which is symmetric with eigenvalues $-\lambda_n \geq \dots \geq -\lambda_1$ and associated eigenvectors v_n, \dots, v_1 , we get

$$v_n = \arg \min_{\|v\|=1} v^\top Av, \quad \text{and for } k = 1, \dots, n-1 \quad v_k = \arg \min_{\|v\|=1, v \perp v_{k+1}, \dots, v_n} v^\top Av.$$

Positive matrices

Definition 1.1

A symmetric matrix $A \in \mathbb{R}^{n \times n}$ is said to be positive semi-definite if

$$\forall x \in \mathbb{R}^n, x^\top A x \geq 0. \quad (1)$$

The matrix A is said to be positive definite if moreover the inequality in (1) is strict for all $x \neq 0$.

Remark 1.2. Negative semi-definite and negative definite matrices are defined analogously.

Proposition 1.2

Let $A \in \mathbb{R}^{n \times n}$ be a symmetric matrix, and let $\lambda_1, \dots, \lambda_n \in \mathbb{R}$ its eigenvalues. Then

$$A \text{ is positive semi-definite} \iff \lambda_i \geq 0 \text{ for } i = 1, \dots, n,$$

and

$$A \text{ is positive definite} \iff \lambda_i > 0 \text{ for } i = 1, \dots, n.$$

Exercise 1.1. Let $A \in \mathbb{R}^{n \times n}$.

- Show that $A^\top A$ positive semi-definite.
- Let M be a $n \times n$ symmetric positive semi-definite matrix. Show that there exists $A \in \mathbb{R}^{n \times n}$ such that $M = A^\top A$.

2 Singular value decomposition

Theorem 2.1 (Singular value decomposition (SVD))

Let $A \in \mathbb{R}^{n \times m}$. Then there exists two orthogonal matrices $U \in \mathbb{R}^{n \times n}$ and $V \in \mathbb{R}^{m \times m}$ and a matrix $\Sigma \in \mathbb{R}^{n \times m}$ such that $\Sigma_{1,1} \geq \Sigma_{2,2} \geq \dots \geq 0$ and $\Sigma_{i,j} = 0$ for $i \neq j$

$$A = U \Sigma V^\top.$$

The columns u_1, \dots, u_n of U (respectively the columns v_1, \dots, v_m of V) are called the left (resp. right) singular vectors of A . The non-negative numbers $\Sigma_{i,i}$ are the singular values of A . Moreover $\text{rank}(A) = \#\{i \mid \Sigma_{i,i} \neq 0\}$.

Theorem 2.1 is proved at the end of these notes.

Notice that the singular vectors (similarly to the eigenvectors) are not uniquely defined: if $A = U \Sigma V^\top$ is a SVD of A , then $A = (-U) \Sigma (-V)^\top$ is also a SVD of A . However, with a slight abuse of language, we will often refer v_i as the i^{th} right singular vector of A .

2.1 Properties of the SVD

Let $A \in \mathbb{R}^{n \times m}$ and let $U \Sigma V^\top$ be a singular value decomposition of A as in Theorem 2.1. Let u_1, \dots, u_n be the left singular vectors (i.e. the columns of U) and v_1, \dots, v_m be the right singular vectors (i.e. the columns of V). Let $\sigma_i = \Sigma_{i,i}$ be the singular values of A .

Proposition 2.1

For $i = 1, \dots, \text{rank}(A)$ we have

$$Av_i = \sigma_i u_i \quad \text{and} \quad A^\top u_i = \sigma_i v_i.$$

The most important property of the singular vectors for us is the following:

Proposition 2.2

We have

$$v_1 = \arg \max_{\|v\|=1} \|Av\| \quad \text{and} \quad \sigma_1 = \max_{\|v\|=1} \|Av\|. \quad (2)$$

It holds also that

$$v_2 = \arg \max_{\|v\|=1, v \perp v_1} \|Av\| \quad \text{and} \quad \sigma_2 = \max_{\|v\|=1, v \perp v_1} \|Av\| \quad (3)$$

and more generally:

$$v_k = \arg \max_{\|v\|=1, v \perp v_1, \dots, v_{k-1}} \|Av\|. \quad \text{and} \quad \sigma_k = \max_{\|v\|=1, v \perp v_1, \dots, v_{k-1}} \|Av\|. \quad (4)$$

Remark 2.1. Considering A^\top leads to an analogous result for the left singular vectors u_k :

$$u_k = \arg \max_{\|u\|=1, u \perp u_1, \dots, u_{k-1}} \|A^\top u\|. \quad \text{and} \quad \sigma_k = \max_{\|u\|=1, u \perp u_1, \dots, u_{k-1}} \|A^\top u\|. \quad (5)$$

Proof. Compute $A^\top A = V \Sigma^\top \Sigma V^\top = V D V^\top$ where the matrix $D \stackrel{\text{def}}{=} \Sigma^\top \Sigma$ is diagonal with $D_{i,i} = \sigma_i^2$. The family (v_1, \dots, v_m) is therefore an orthonormal family of eigenvectors of the symmetric matrix $A^\top A$ and $\sigma_1^2 \geq \dots \geq \sigma_m^2$ are the corresponding eigenvalues. The result follows then from Proposition 1.1 applied to $A^\top A$, noticing that $v^\top A^\top A v = \|Av\|^2$. \square

3 Interpretation and applications of the SVD

3.1 Geometric interpretation

The decomposition $M = U \Sigma V^\top$ gives that the linear transformation associated to the matrix M is the composition of three linear transformations:

1. V^\top is a rotation/reflection: length, dot products, angles are preserved.
2. Σ corresponds to a scaling.
3. U is another rotation/reflection: length, dot products, angles are preserved.




Figure 1: Geometric interpretation of SVD

3.2 “Maximal variance” interpretation

Let $a_1, \dots, a_n \in \mathbb{R}^d$ be n points in d dimensions. We assume that these points are centered, meaning that

$$\sum_{i=1}^n a_i = 0.$$

Let A be the $n \times d$ matrix whose rows are a_1, \dots, a_n and let (v_1, \dots, v_n) be its right singular vectors. By Proposition 2.2, v_1 , the first right singular vector of A , maximizes

$$v \mapsto \|Av\|^2 = \sum_{i=1}^n \langle a_i, v \rangle^2$$

over the unit sphere $\mathbb{S}_n \stackrel{\text{def}}{=} \{x \in \mathbb{R}^n \mid \|x\| = 1\}$. This quantity is the variance of the coordinates of the points a_1, \dots, a_n along the direction $\text{Span}(v)$.

The first right singular vector v_1 gives therefore the direction along which the variance of the data is maximal. Proposition 2.2 gives also that

$$v_k = \arg \max_{\|v\|=1, v \perp v_1, \dots, v_{k-1}} \|Av\|^2. \quad (6)$$

Hence v_2 gives the direction orthogonal to v_1 that maximizes the variance and so on...

3.3 Application: Principal Component Analysis (PCA)

The SVD is commonly used for dimensionality reduction. Assume that we are given a dataset of n points $a_1, \dots, a_n \in \mathbb{R}^d$, with d very large. We aim at representing this dataset in lower dimension, i.e. finding $\tilde{a}_1, \dots, \tilde{a}_n \in \mathbb{R}^k$ where k is smaller than d , such that the points $(\tilde{a}_1, \dots, \tilde{a}_n)$ look like the original ones (a_1, \dots, a_n) . This could be for instance used

- to reduce computing time.
- to visualize an high-dimensional dataset in dimension $k = 2$ or 3 .

Assume that the dataset is centered (otherwise subtract the mean $\mu = \frac{1}{n} \sum_{i=1}^n a_i$ to all the points). As we have seen in Section 3.2 above, the k first right-singular vectors v_1, \dots, v_k gives the k first orthogonal direction along which the variance of the dataset is maximal.

Hence, it is very natural to project the original dataset a_1, \dots, a_k on these directions $\text{Span}(v_1, \dots, v_k)$: we are going to represent each point a_i by the coordinates of its projection, in the basis v_1, \dots, v_k . Since (v_1, \dots, v_k) is orthonormal, the coordinate of a_i along the vector v_j is $\langle a_i, v_j \rangle$. Consequently we take

$$\tilde{a}_i = (\langle v_1, a_i \rangle, \dots, \langle v_k, a_i \rangle) = V_{|k}^T a_i,$$

where $V_{|k} = (v_1 | \dots | v_k) \in \mathbb{R}^{d \times k}$.

How do we chose k ? The dimension k of the output space can be chosen by looking at the singular values of A . Let $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_{\min(n,d)}$ be the singular values of A . As seen in Section 3.2, the variance of the dataset along the direction v_i is σ_i^2 .




Figure 2: Singular values of A , ranked in decreasing order.

A simple way to chose k is therefore to plot the square singular values as on Figure 2 and look for a good cut-off ($k = 4$ on Figure 2). Doing so, one captures a fraction

$$\frac{\sum_{i=1}^k \sigma_i^2}{\sum_{i=1}^{\min(n,d)} \sigma_i^2}$$

of the total variance.

Should we “normalize” the dataset? It depends, but in general the answer is yes, especially if you have data from heterogeneous types. Imagine that you have measured the size and the weight of n objects, and stored the information in vectors $a_i = (\text{size of object } i \text{ in cm}, \text{weight of object } i \text{ in kg})$. If I change the weighting unit from kilograms to grams, this multiply the variance along the second coordinates by 10^6 , leading to very different principal components. Normalizing the dataset (i.e. dividing the columns of the data matrix by their standard deviation) allows to be unaffected by a change of units.

However, this decreases a lot the variance of columns which high variance and amplify a lot the variance of columns with low variance. Hence you may not always want to normalize the columns.

3.4 Best-fitting subspace

Let $a_1, \dots, a_n \in \mathbb{R}^d$ be n points in d dimensions. We consider the problem of finding the k -dimensional subspace (for $k = 1, \dots, n$) that fits “the best” these n data points. By “best”, we mean here the k -dimensional subspace S that minimize the sum of the square distances to the n points:

$$\text{minimize } \sum_{i=1}^n d(a_i, S)^2 \text{ with respect to } S \text{ subspace of dimension } k. \quad (7)$$

Recall that the distance of a vector x to the subspace S is defined as $d(x, S) = \|x - P_S x\|$. Let A be the $n \times d$ matrix whose rows are a_1, \dots, a_n . The goal of this section is to prove:

Theorem 3.1

Let v_1, \dots, v_n be right singular vectors of A . Then for all $k \in \{1, \dots, n\}$, the subspace $\text{Span}(v_1, \dots, v_k)$ is a solution of (7).

We start by noticing that for all $i \in \{1, \dots, n\}$,

$$d(a_i, S)^2 = \|a_i - P_S(a_i)\|^2 = \|a_i\|^2 - \|P_S(a_i)\|^2,$$

by Pythagorean Theorem (recall that $P_S(a_i) \perp (a_i - P_S(a_i))$). Minimizing (7) is therefore equivalent to maximize

$$\sum_{i=1}^n \|P_S(a_i)\|^2. \quad (8)$$

Let us fix an orthonormal basis (s_1, \dots, s_k) of S . Then for all $x \in \mathbb{R}^d$, $P_S(x) = \langle s_1, x \rangle s_1 + \dots + \langle s_k, x \rangle s_k$, hence

$$\sum_{i=1}^n \|P_S(a_i)\|^2 = \sum_{i=1}^n \sum_{j=1}^k \langle a_i, s_j \rangle^2 = \|As_1\|^2 + \dots + \|As_k\|^2. \quad (9)$$

Consequently, minimizing (7) is equivalent to maximizing (9) over all orthonormal families (s_1, \dots, s_k) .

For simplicity, we start by considering the case $k = 1$, in which case $S = \text{Span}(s_1)$. In this case, (9) is simply:

$$\sum_{i=1}^n \|P_S(a_i)\|^2 = \|As_1\|^2. \quad (10)$$

Proposition 2.2 tells us that a subspace of dimension 1 that maximizes (10) and hence that minimizes (7) is $\text{Span}(v_1)$ because

$$v_1 = \arg \max_{\|v\|=1} \|Av\|. \quad (11)$$

If we now want to solve the problem for $k = 2$, a natural candidate for the subspace S would be $S = \text{Span}(v_1, v_2)$ since by Proposition 2.2

$$v_2 = \arg \max_{\|v\|=1, v \perp v_1} \|Av\|. \quad (12)$$

We can follow this greedy strategy for $k = 3, \dots, n$, $S = \text{Span}(v_1, \dots, v_k)$ is a natural candidate for being solution of (7).

It is not a priori obvious (except for $k = 1$) that $S = \text{Span}(v_1, \dots, v_k)$ is a minimizer of (7) over all the subspaces of dimension k . We need the following lemma.

Lemma 3.1

Let $k \in \{2, \dots, k\}$. Assume that (v_1, \dots, v_{k-1}) is an orthonormal family that maximizes (9). Define

$$v_k = \arg \max_{\|v\|=1, v \perp \text{Span}(v_1, \dots, v_{k-1})} \|Av\|.$$

Then (v_1, \dots, v_k) is an orthonormal family and $\text{Span}(v_1, \dots, v_k)$ minimizes (7), i.e. (v_1, \dots, v_k) maximizes (9).

Proof. Let S be a subspace of dimension k . Let (w_1, \dots, w_k) be an orthonormal basis of S such that $w_k \perp \text{Span}(v_1, \dots, v_{k-1})$. By definition of v_k , we have $\|Aw_k\| \leq \|Av_k\|$. We also assumed that (v_1, \dots, v_k) maximizes (9), so

$$\|Av_1\|^2 + \dots + \|Av_{k-1}\|^2 \geq \|Aw_1\|^2 + \dots + \|Aw_{k-1}\|^2.$$

We conclude that

$$\|Av_1\|^2 + \dots + \|Av_k\|^2 \geq \|Aw_1\|^2 + \dots + \|Aw_k\|^2,$$

so (v_1, \dots, v_k) maximizes (9). \square

Theorem 3.1 follows then by induction.

Proof of Theorem 2.1

We apply the Spectral Theorem (Theorem 1.1) to the $m \times m$ matrix $A^\top A$: there exists an orthonormal basis (v_1, \dots, v_m) of \mathbb{R}^m of eigenvectors of $A^\top A$ associated to eigenvalues $\lambda_1 \geq \dots \geq \lambda_m$ that are all non-negative because $A^\top A$ is non-negative. Let $V \in \mathbb{R}^{m \times m}$ be the orthogonal matrix whose columns are (v_1, \dots, v_m) .

Let us write $\sigma_i = \sqrt{\lambda_i}$ and let $r = \max\{i | \sigma_i > 0\}$. Define for $i = 1, \dots, r$

$$u_i = \frac{1}{\sigma_i} Av_i \in \mathbb{R}^n. \quad (13)$$

Lemma 3.2

The family (u_1, \dots, u_r) is orthonormal.

Proof. Let $i, j \in \{1, \dots, r\}$.

$$\langle u_i, u_j \rangle = \left(\frac{1}{\sigma_i} Av_i \right)^\top \left(\frac{1}{\sigma_j} Av_j \right) = \frac{1}{\sigma_i \sigma_j} v_i^\top A^\top A v_j = \frac{\sigma_i}{\sigma_j} v_i^\top v_j = \mathbb{1}_{i=j},$$

since $A^\top A v_i = \sigma_i^2 v_i$. \square

If $r < n$ we let (u_{r+1}, \dots, u_n) be an orthonormal family of vectors of \mathbb{R}^n that are orthogonal to u_1, \dots, u_r . The family (u_1, \dots, u_n) is then an orthonormal basis of \mathbb{R}^n . Let $U \in \mathbb{R}^{n \times n}$ be the orthogonal matrix whose columns are (u_1, \dots, u_n) .

Lemma 3.3

For $i = r+1, \dots, m$, $Av_i = 0$.

Proof. We compute for $i = r+1, \dots, m$:

$$\|Av_i\|^2 = v_i^\top A^\top A v_i = v_i^\top (\lambda_i v_i) = \sigma_i^2 = 0.$$

□

Finally, we let $\Sigma \in \mathbb{R}^{n \times m}$ defined by:

$$\Sigma_{i,j} = \begin{cases} \sigma_i & \text{if } i = j \\ 0 & \text{otherwise.} \end{cases}$$

It remains to verify that $A = U\Sigma V^\top$. Compute for $i = 1, \dots, m$, using the definition (13) and Lemma 3.3:

$$Av_i = \begin{cases} \sigma_i u_i & \text{if } i \leq r \\ 0 & \text{otherwise.} \end{cases}$$

By orthogonality of V and the construction of Σ one verifies easily that

$$U\Sigma V^\top v_i = \begin{cases} \sigma_i u_i & \text{if } i \leq r \\ 0 & \text{otherwise.} \end{cases}$$

We conclude that for all $i \in \{1, \dots, m\}$, $Av_i = U\Sigma V^\top v_i$. Since a linear transformation is uniquely determined by the image of a basis, we conclude that $A = U\Sigma V^\top$.

It remains to show:

Lemma 3.4

rank(A) = r .

Proof. The family (u_1, \dots, u_r) is orthonormal, hence linearly independent. By definition $u_i \in \text{Im}(A)$ which implies that $\text{rank}(A) = \dim(\text{Im}(A)) \geq r$. To prove the converse inequality, notice that by Lemma 3.3 $v_i \in \text{Ker}(A)$ for $i = r+1, \dots, m$. The vectors (v_{r+1}, \dots, v_m) are orthonormal, hence linearly independent. This implies that $\dim(\text{Ker}(A)) \geq m - r$. We conclude by applying the rank Theorem:

$$\text{rank}(A) = m - \dim(\text{Ker}(A)) \leq m - (m - r) = r.$$

□

